

15-319 / 15-619

Cloud Computing

Recitation 5

February 14th, 2017

Overview

Administrative issues

Office Hours, Piazza guidelines

Last week's reflection

Project 2.1, OLI Unit 2 modules 5 and 6

This week's schedule

Quiz 4 - February 17, 2017 (Modules 7, 8, 9)

Project 2.2 - February 19, 2017

Finalize 3 person teams for the Team Project -
February 20, 2017

Announcements



- Monitor your expenses regularly
 - Check your bill frequently on TheProject.Zone
 - Check on AWS, use Cost Explorer & filter by tags
 - Check on the Azure portal since only \$100/sem
- **Terminate** your resources when you are done
 - Stopping a VM still costs EBS money (\$0.1/GB/month)
 - Amazon EC2 and Amazon Cloudwatch fees for monitoring, ELB
 - AutoScaling Group - no additional fees

Announcements



Use spot instances as much as possible

Complete the Project Survey for P1.1, P1.2, P2.1
Piazza @1216

Protect your credentials

Crawlers are looking for AWS credentials on
public repos!

Primer for 3.1 is out

Storage I/O Benchmarking

Last Week's Reflection



OLI: Conceptual Content

Unit 2 - Modules 5 and 6:

Cloud Management & Software Deployment
Considerations

Quiz 3 completed

P2.1: Azure, GCP, and AWS EC2 APIs

CLI, Java, Python

P2.1: Load Balancing and AutoScaling

Experience horizontal scaling

Programmatically manage cloud resources and
deal with failure

Initial experience with load balancing

Project 2.1



To evaluate how well other people can read your code, we will be manually grading your submitted code

Azure

GCP

AWS (Horizontal and Autoscaling)

To enhance readability

Use the [Google Code Style](#) guidelines

Always add comments especially for complex parts

Project 2.1




You gained experience working with Cloud Service Provider's APIs

Documentation may be unclear or may omit certain details.

Have to ensure you're referencing the correct documentation version.

Considerations of cost vs performance

This Week: Content



- UNIT 3: Virtualizing Resources for the Cloud
 - Module 7: Introduction and Motivation
 - Module 8: Virtualization
 - Module 9: Resource Virtualization - CPU
 - Module 10: Resource Virtualization - Memory
 - Module 11: Resource Virtualization – I/O
 - Module 12: Case Study
 - Module 13: Network and Storage Virtualization

OLI Module 7 - Virtualization

Introduction and Motivation

Why virtualization

Enabling the cloud computing system model

Elasticity

Resource sandboxing

Limitation of General-Purpose OS


Mixed OS environment

Resource sharing

Time

Space

Improved system utilization and reduce costs
from a cloud provider perspective



OLI Module 8 - Virtualization

What is Virtualization

Involves the construction of an isomorphism that maps a virtual guest system to a real (or physical) host system


Sequence of operations e modify guest state

Mapping function $V(S_i)$

Virtual Machine Types

Process Virtual Machines

System Virtual Machines



OLI Module 9

Resource Virtualization - CPU

Steps of CPU Virtualization

- Multiplexing a physical CPU among virtual CPUs


- Virtualizing the ISA (Instruction Set Architecture) of a CPU

- Code Patch, Full Virtualization and Paravirtualization

- Emulation (Interpretation & Binary Translation)

- Virtual CPU

This Week: Project



P2.1: Horizontal Scaling and Autoscaling
MSB Interview

P2.2: Containers and Kubernetes
Building a Coding Interview Playground
Working with Docker and Kubernetes

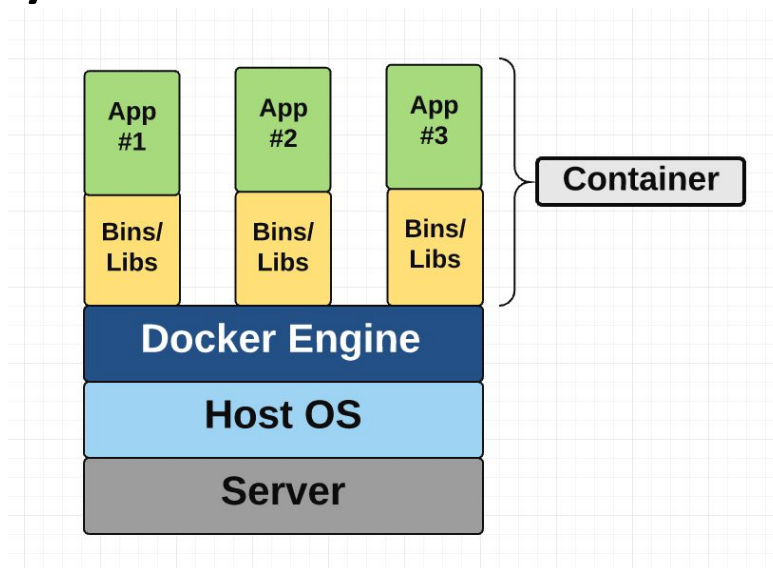
Containers



Provides OS-level virtualization.

Provides private namespace, network interface and IP address, etc.

Big difference with VMs is that containers share the host system's kernel with other containers.



Why Containers?



Faster deployment


Portability across machines

Version control

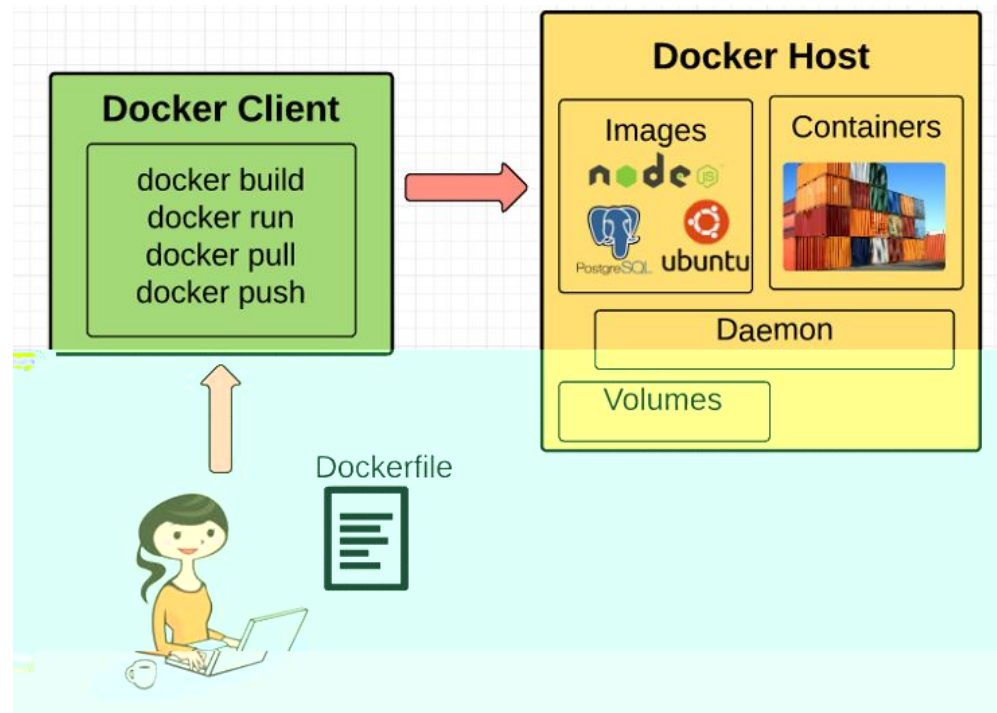
Simplified dependency management

Build once, run anywhere


Docker Engine



An orchestrator that comprises:
Docker Daemon
Docker Client
REST API



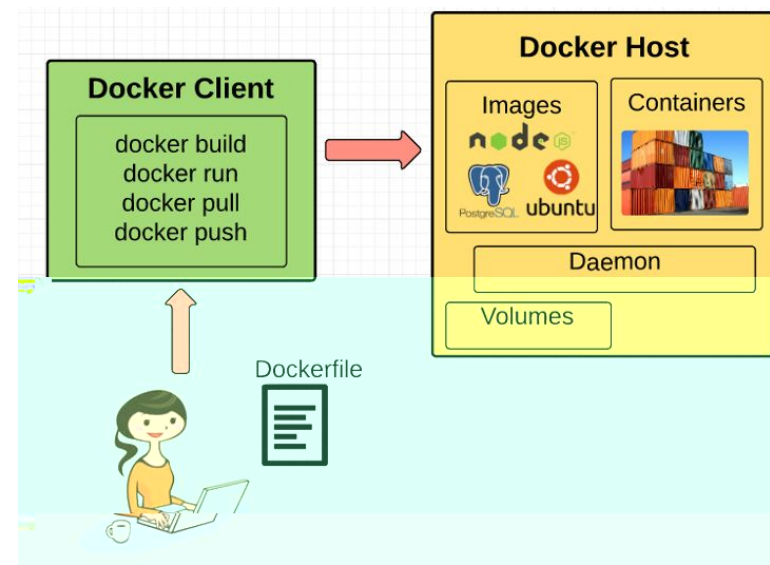
Docker Daemon




Main brains behind Docker Engine

The Docker Client is used to communicate with the Docker Daemon

The Daemon does not have to be on the same machine as the Client



Docker Client

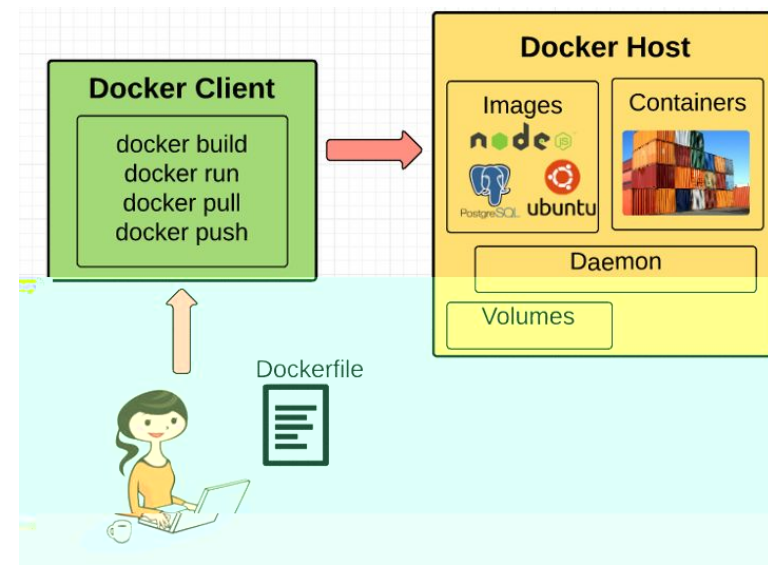


Also known as Docker CLI


When you type:

```
docker build nginx
```

You are telling the Docker client to forward the
build nginx
instruction to the Daemon



Dockerfile



We can use a Dockerfile to build container images

Dockerfile tells Docker:

- What base image to work from

- What commands to run on base image

- What files to copy to the base image

- What ports should the container listen on?

- What binaries to execute when the container launches?

In short, Dockerfile is a recipe for Docker images

Let's go through a sample Dockerfile!

Example Dockerfile


```
# Debian as the base image  
FROM debian:latest
```

```
# Install additional packages  
RUN apk add --update emacs  
RUN apk add --update apache
```

```
# index.html must be in the current directory  
ADD index.html /home/demo/
```

```
# Define the command which runs when the container starts  
CMD ["cat /home/demo/index.html"]
```

```
# Use bash as the container's entry point. CMD is the argument to this entry  
point  
ENTRYPOINT ["/bin/bash", "-c"]
```



Example Dockerfile


```
# Debian Linux as the base image  
FROM debian:latest
```

```
# Install additional packages  
RUN apk add --update emacs  
RUN apk add --update apache
```

```
# index.html must be in the current directory  
ADD index.html /home/demo/
```

```
# Define the command which runs when the container starts  
CMD ["cat /home/demo/index.html"]
```

```
# Use bash as the container's entry point. CMD is the argument to this entry  
point  
ENTRYPOINT ["/bin/bash", "-c"]
```



Example Dockerfile

```
# Alpine Linux as the base image
```

```
FROM debian:latest
```

```
# Install additional packages
```

```
RUN apk add --update emacs
```

```
RUN apk add --update apache
```

```
# index.html must be in the current directory
```


```
ADD index.html /home/demo/
```

```
# Define the command which runs when the container starts
```

```
CMD ["cat /home/demo/index.html"]
```

```
# Use bash as the container's entry point. CMD is the argument to this entry point
```

```
ENTRYPOINT ["/bin/bash", "-c"]
```



Example Dockerfile


```
# Alpine Linux as the base image  
FROM debian:latest
```

```
# Install additional packages  
RUN apk add --update emacs  
RUN apk add --update apache
```

```
# index.html must be in the current directory  
ADD index.html /home/demo/
```

```
# Define the command which runs when the container starts  
CMD ["cat /home/demo/index.html"]
```

```
# Use bash as the container's entry point. CMD is the argument to this entry  
point  
ENTRYPOINT ["/bin/bash", "-c"]
```



Example Dockerfile

```
# Alpine Linux as the base image  
FROM debian:latest
```


```
# Install additional packages  
RUN apk add --update emacs  
RUN apk add --update apache
```

```
# index.html must be in the current directory  
ADD index.html /home/demo/
```

```
# Define the command which runs when the container starts  
CMD ["cat /home/demo/index.html"]
```

```
# Use bash as the container's entry point. CMD is the argument to this entry  
point
```

```
ENTRYPOINT ["/bin/bash", "-c"]
```



Images & Containers

An image: is a static file; never changes

A container: a live instance of an image

Think of it this way – you have a DVD that installs Windows OS (image). After you install it, you can write files to it (container).

`docker build`

builds an image

`docker run`

runs a container based off of an image

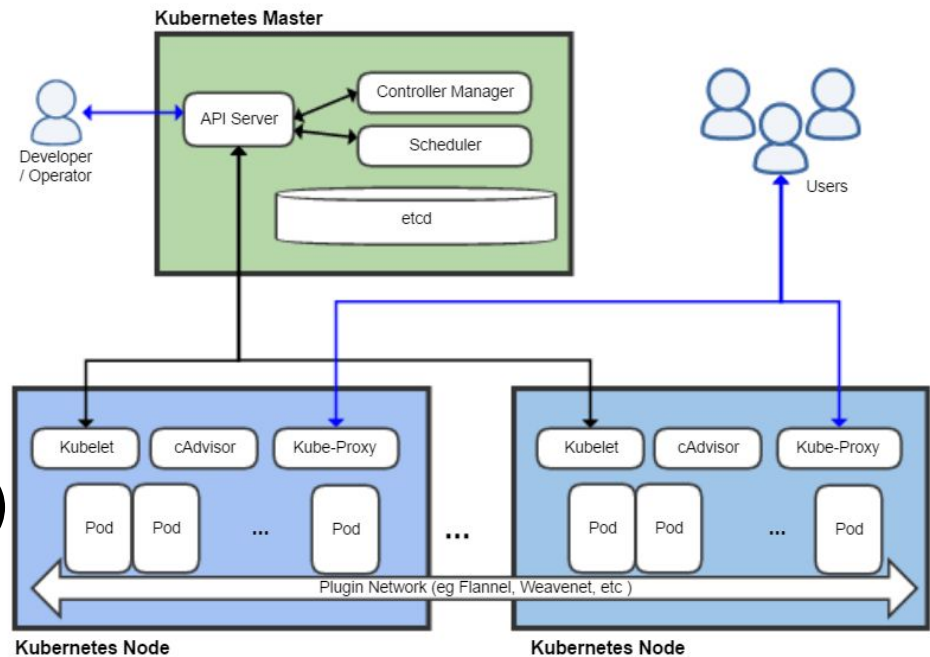
Kubernetes

Kubernetes is an open-source platform for automating deployment, scaling, and operations of application containers.

Containers built using Docker

Images stored in a private registry

Application defined using a YAML (or JSON) template



Deploying with Kubernetes


[kubect](#)l is the command that interacts with the
Kubernetes Master (API Server)

Pod Configuration [Reference](#)

```
kubect
```

```
l create -f demo_pod.yaml

apiVersion: v1
kind: Pod
metadata:
  name: test
  labels:
    app: test
spec:
  containers:
    - name: backend1
      image: us.gcr.io/cc-p22/task2-container1:v0
      ports:
        - name: http
          containerPort: 8080
  ...
```



Project 2.2 - Containers

Build a service to compile and run user code submitted through a front end.

Four tasks:

Task 1: Containerize an Nginx server and run container locally.

Task 2: Build a service to call two web services and return the sum. Deploy application to a Kubernetes cluster.

Task 3: Build a Python web service to evaluate Python code submitted from the UI and return the result.

Task 4: Multi-cloud deployment. Python and SQL code evaluation services. Multiple application + DB in the same pod.

Task 1 Objectives

Work with Dockerfiles

Master the Docker CLI, including useful commands like:

- `docker build`
- `docker images`
- `docker run`
- `docker ps`


Think about integration between the host and the container

Task 1 Overview

Configure a Docker container with an Nginx web server

Nginx server listening on port 15619

Port 15619 of host VM mapped to the container port



Task 2 Objectives


Work with more complex Dockerfiles

Define a Kubernetes YAML file for your application

Deploying your application to the cluster

Multiple pods per container

Exposing the pod as a Service via ELB



Task 3 Objectives

Developing a code execution service

- Accepts Python code from the UI (over HTTP)

- Use the provided API specification to develop your application.

 - Service that consumes and produces JSON.

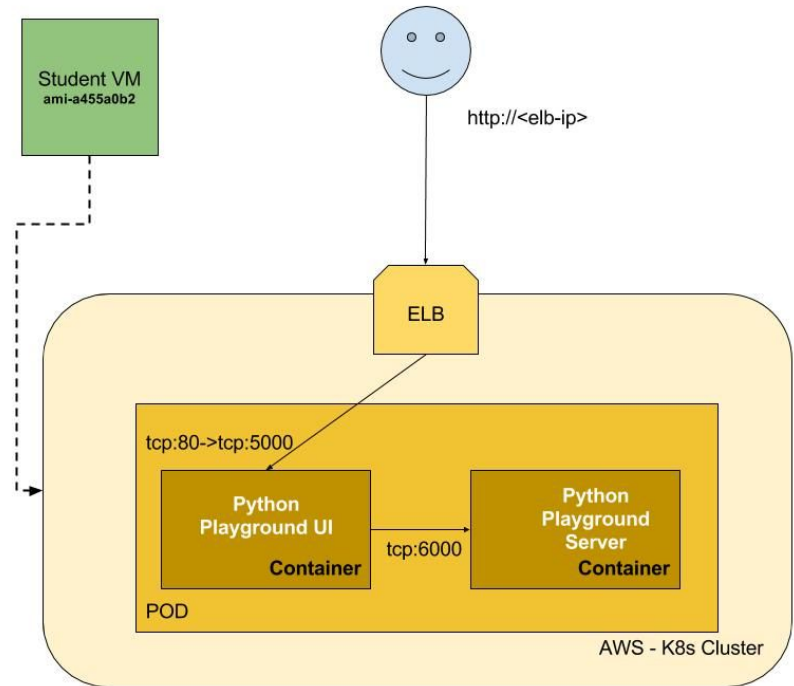
Develop Kubernetes YAML definition


- Application for UI, exposed as Service via LB

- Backend service that is not exposed to the Internet, only accessible to the UI container

Task 3 Overview

Python code evaluation service architecture.
The UI is exposed on the internet, accepts POST requests.



 Python

Python Code

```
1 # Welcome to the Python Explorer
2 import random;
3
4 firstRandom = random.randint(1, 100);
5 secondRandom = random.randint(1, 100);
6
7 str(firstRandom);
8 str(secondRandom);
9
10 : " + str(sum));
```

Output

```
1 The first random number is: 34
2 The second random number is: 74
3 The sum of the random numbers is: 108
4
```

```
6
7 sum = firstRandom + secondRandom;
8
9 print ("The first random number is: " +
10 print ("The second random number is: "
11
12 print ("The sum of the random numbers i
13
```

Submit

Task 4 Objectives

Task 4 will build on task 3 by introducing a SQL evaluation service.

Deploy a Java REST service and MySQL database to your Kubernetes cluster.

Achieve high availability!

Multi cloud deployment!

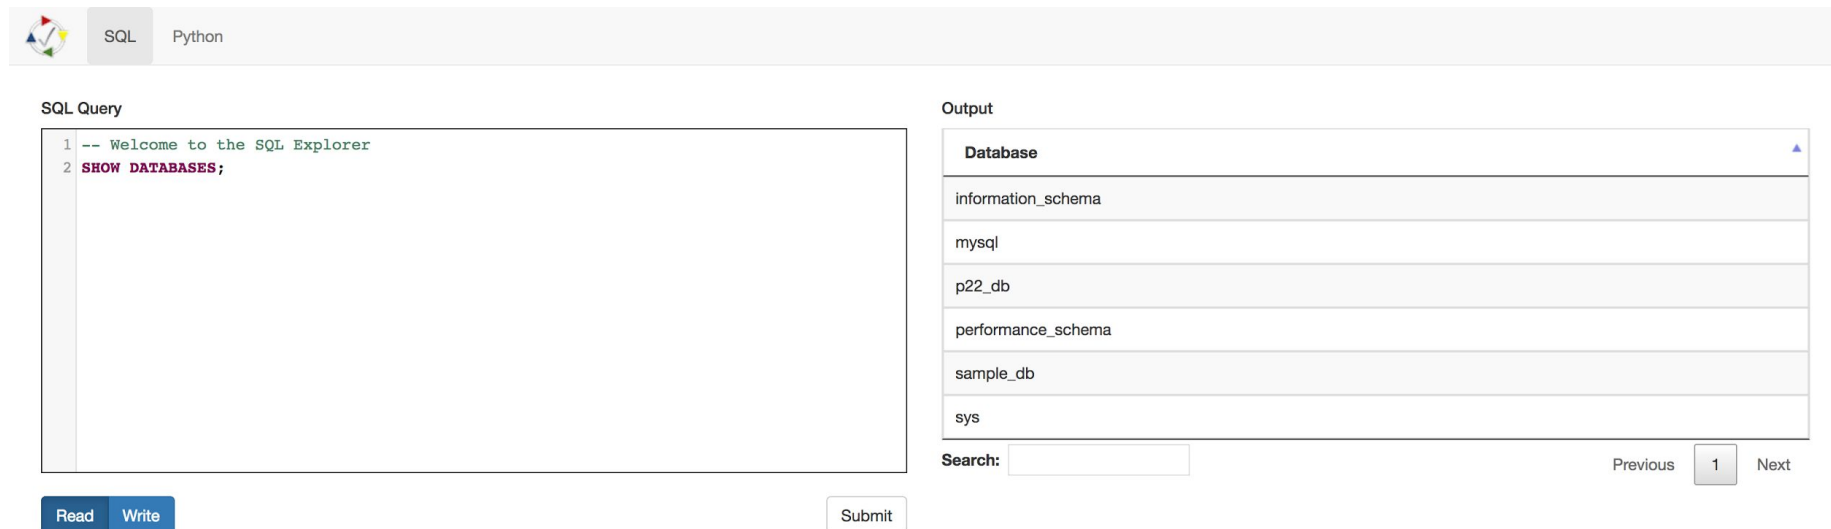
The Python evaluation service will be replicated across multiple clouds.

Task 4 Overview

Multi-cloud application deployments.

High availability! Python evaluation service in GCP and Azure.

UI will handle SQL and Python code samples.



The screenshot displays a web-based SQL Explorer interface. At the top, there is a navigation bar with a logo on the left and two tabs: 'SQL' (which is active) and 'Python'. Below the navigation bar, the interface is split into two main sections. The left section, titled 'SQL Query', contains a text area with two lines of SQL code: a comment '-- Welcome to the SQL Explorer' and the command 'SHOW DATABASES;'. Below the text area are two buttons, 'Read' and 'Write'. The right section, titled 'Output', displays the results of the query in a table format. The table has a single column labeled 'Database' and lists several databases: 'information_schema', 'mysql', 'p22_db', 'performance_schema', 'sample_db', and 'sys'. Below the table is a search bar with the label 'Search:'. At the bottom right of the output section, there are navigation controls: a 'Previous' button, a page indicator '1' inside a box, and a 'Next' button. A 'Submit' button is also located at the bottom center of the interface.

SQL Explorer

SQL Python

SQL Query

```
1 -- Welcome to the SQL Explorer
2 SHOW DATABASES;
```

Read Write

Submit


Output

Database
information_schema
mysql
p22_db
performance_schema
sample_db
sys

Search:

Previous 1 Next

Task 4 Architecture



Tips, Trips, and Tricks

Debug, debug, debug

This project has many moving pieces!

Where is the issue occurring?

What is the expected behaviour of the system?

Pods and Logs

Did my pod start? `(kubectl get pods , kubectl describe pods)`

Is my pod generating any logs? `(kubectl logs ...)`

Project 2.2 Penalties


Danger

Project Grading Penalties

Violation	Penalty of the project grade
Spending more than \$5 for this project on AWS	-10%
Spending more than \$10 for this project on AWS	-100%
Failing to tag all your resources in any task (EC2 instances, ELB) for this project; Key: Project and Value: 2.2	-10%

Submitting your AWO, OOD, or A-ware credentials, either manually or through ID in

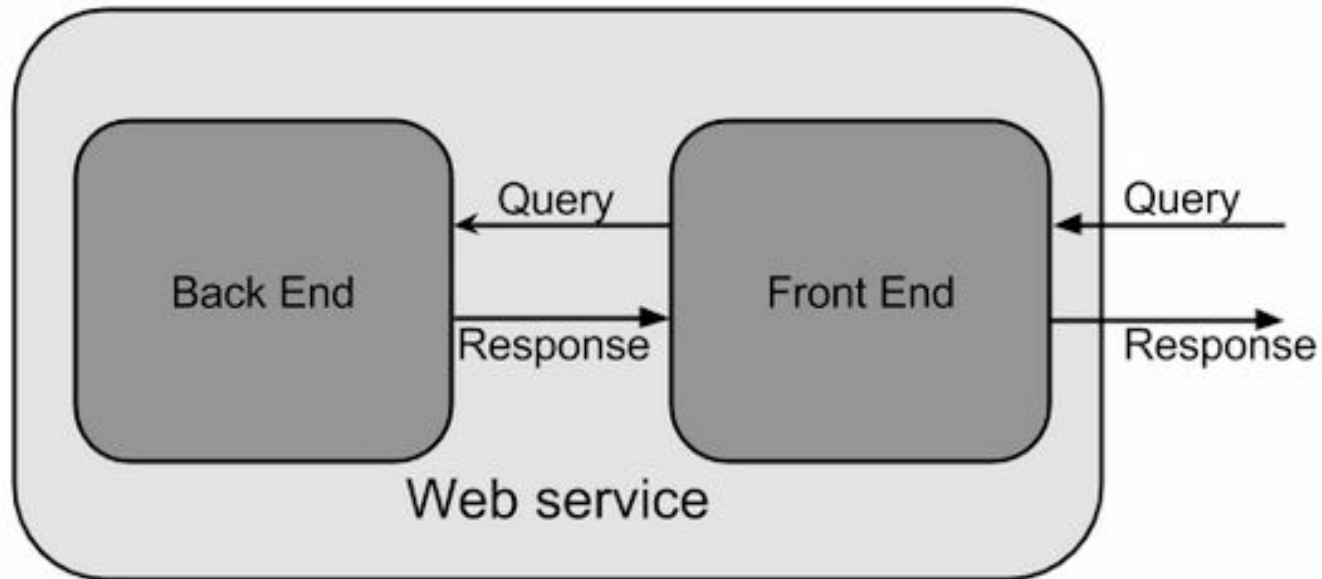
Upcoming Deadlines



- Quiz 4: Modules 7, 8 and 9:
Due: **Friday February 17, 2017 11:59PM Pittsburgh**
- Project 2.2: Docker and Kubernetes
Due: **Sunday February, 19 2017 11:59PM Pittsburgh**
- Team Project: Team Formation
Due: **February, 20 2017 11:59PM Pittsburgh**



Team Project Architecture




Writeup and Queries will be released on Monday, **February 27th, 2017**

We can have more discussions in subsequent recitations

For now, ensure 3-person teams you decide have experience with web frameworks and database, storage principles and infra setup/hacking

Team Project Time Table



Phase (and query due)	Start	Deadline	Code and Report Due
Phase 1 Q1, Q2, Q3	Monday 02/27/2017 00:00:00 ET	Sunday 03/26/2017 23:59:59 ET	Tuesday 03/28/2017 23:59:59 ET
Phase 2 Q1, Q2, Q3, Q4	Monday 03/27/2017 00:00:00 ET	Sunday 04/16/2017 15:59:59 ET	
Phase 2 Live Test (Hbase/MySQL) Q1, Q2, Q3, Q4	Sunday 04/16/2017 18:00:00 ET	Sunday 04/16/2017 23:59:59 ET	Tuesday 04/18/2017 23:59:59 ET
Phase 3 Q1, Q2, Q3, Q4, Q5	Monday 04/17/2017 00:00:00 ET	Sunday 04/30/2017 15:59:59 ET	
Phase 3 Live Test Q1, Q2, Q3, Q4, Q5	Sunday 04/30/2017 18:00:00 ET	Sunday 04/30/2017 23:59:59 ET	Tuesday 05/02/2017 23:59:59 ET

Demo

Configure simple containerized application

Deploy to kubernetes cluster